

# PROJECT REPORT

## CMPS 470/570

Team: Tech19

Team members: Drew Hutchinson, Zichuo Wang

Instructor: Dr. Ömer Soysal

# Report structure

- Description of the project
- Description of the raw data
- Preprocessing
- Feature extraction
- Description of the feature data
- Description of the model
- Performance of the model
- Team members & roles / Task completion report

# Description of the project

Develop a machine learning application using Python and related libraries, to solve the **penguin species classification** problem with the dataset **Palmer Penguins**.

This application can:

Pre-process data,

Create ML models based on the methods of K-NN, ANN, SVM and DT.

Display and export results and model parameters.

# Description of the raw data

Data source: Palmer Penguins

<https://github.com/allisonhorst/palmerpenguins>

## Attributes and their types

Attributes	<b>SID</b>	<b>species</b>	<b>island</b>	<b>bill_length_mm</b>	<b>bill_depth_mm</b>	<b>flipper_length_mm</b>	<b>body_mass_g</b>	<b>sex</b>	<b>year</b>
Types	Numerical	Texts	Texts	Numerical	Numerical	Numerical	Numerical	Texts	Numerical
Sample	2	Adelie	Torgersen	39.5	17.4	186	3800	female	2007

[cmpsML\\_Tech19\OTHER\penguins.csv](#)

[cmpsML\\_Tech19\OTHER\data\\_description.xlsx\raw\\_data\\_description](#)

Discussion:

We downloaded the file penguins.csv as our original dataset.

It has 9 columns and 344 rows. The column ‘species’ is the target column.

# Preprocessing

We imported the original data set from penguins.csv.

We defined two functions handling the missing values.

We also applied scalers to the dropped missing values data set.

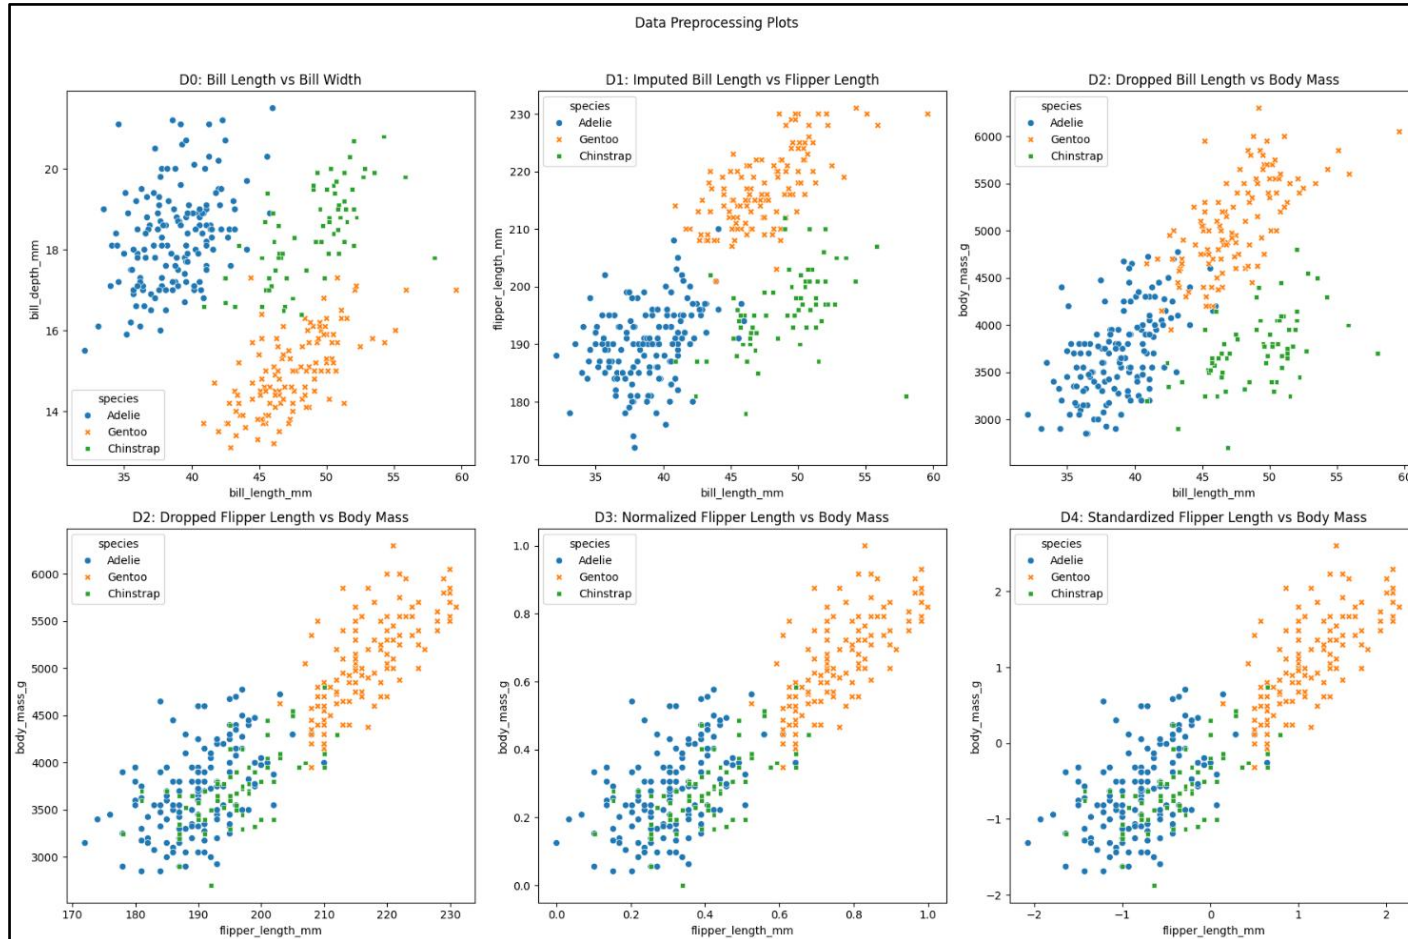
```
"<preprocessing>" module begins.  
Original data set D0:  
This Data Set has 344 rows, 9 columns  
  SID species    island ... body_mass_g    sex    year  
0    1 Adelie  Torgersen ...    3750.0    male  2007  
1    2 Adelie  Torgersen ...    3800.0  female  2007  
2    3 Adelie  Torgersen ...    3250.0  female  2007  
3    4 Adelie  Torgersen ...      NaN     NaN   2007  
4    5 Adelie  Torgersen ...    3450.0  female  2007
```

```
Imputed missing values data set D1:  
This Data Set has 344 rows, 9 columns  
  SID species    island ... body_mass_g    sex    year  
0    1 Adelie  Torgersen ...  3750.000000    male  2007  
1    2 Adelie  Torgersen ...  3800.000000  female  2007  
2    3 Adelie  Torgersen ...  3250.000000  female  2007  
3    4 Adelie  Torgersen ...  4201.754386    male  2007  
4    5 Adelie  Torgersen ...  3450.000000  female  2007  
  
[5 rows x 9 columns]  
Dropped missing values data set D2:  
This Data Set has 333 rows, 9 columns  
  SID species    island ... body_mass_g    sex    year  
0    1 Adelie  Torgersen ...    3750.0    male  2007  
1    2 Adelie  Torgersen ...    3800.0  female  2007  
2    3 Adelie  Torgersen ...    3250.0  female  2007  
4    5 Adelie  Torgersen ...    3450.0  female  2007  
5    6 Adelie  Torgersen ...    3650.0    male  2007
```

```
Normalized data set D3:  
This Data Set has 333 rows, 9 columns  
  SID species    island ... body_mass_g    sex    year  
0    1 Adelie  Torgersen ...    0.291667    male  2007  
1    2 Adelie  Torgersen ...    0.305556  female  2007  
2    3 Adelie  Torgersen ...    0.152778  female  2007  
4    5 Adelie  Torgersen ...    0.208333  female  2007  
5    6 Adelie  Torgersen ...    0.263889    male  2007  
  
[5 rows x 9 columns]  
Standardized data set D4:  
This Data Set has 333 rows, 9 columns  
  SID species    island ... body_mass_g    sex    year  
0    1 Adelie  Torgersen ...   -0.568475    male  2007  
1    2 Adelie  Torgersen ...   -0.506286  female  2007  
2    3 Adelie  Torgersen ...   -1.190361  female  2007  
4    5 Adelie  Torgersen ...   -0.941606  female  2007  
5    6 Adelie  Torgersen ...   -0.692852    male  2007
```

# Preprocessing

Here is our data preprocessing plots:



Discussion:

In this part we tried to read files from the folder, apply scalers to the dataset and generate plots.

In the second row of the plots, we compared flipper length vs body mass from 3 datasets: Dropped missing value dataset D2, Normalized dataset D3, Standardized dataset D4.

You can see that **scaler won't change the distribution of the data.**

# Feature extraction

We dropped the rows where have missing values.

We dropped the irrelevant column('SID' , 'Year').

We applied one-hot encoding for categorical columns('sex', 'island') and label encoding for species column.

```
def preprocessing_penguins(df): 1 usage
    df = df.copy()
    df = drop_missing_values(df)
    df = drop_irrelevant_columns(df)
    df = one_hot_encode(df)
    df = label_encode_species(df)
    return df
```

```
"<FeatureExtraction>" module begins.
Original dataset D0:
This Data Set has 344 rows, 9 columns
  SID species  island  ...  body_mass_g  sex  year
0   1  Adelie  Torgersen  ...    3750.0  male  2007
1   2  Adelie  Torgersen  ...    3800.0  female  2007
2   3  Adelie  Torgersen  ...    3250.0  female  2007
3   4  Adelie  Torgersen  ...         NaN    NaN  2007
4   5  Adelie  Torgersen  ...    3450.0  female  2007

[5 rows x 9 columns]
Preprocessed dataset D1:
This Data Set has 333 rows, 9 columns
  species  bill_length_mm  ...  island_Dream  island_Torgersen
0       0             39.1  ...           0                1
1       0             39.5  ...           0                1
2       0             40.3  ...           0                1
4       0             36.7  ...           0                1
5       0             39.3  ...           0                1

[5 rows x 9 columns]
```



# Feature extraction

We divided the dataset into 3 sub-sets ( train, val, test = 0.6, 0.2, 0.2).

We added some noise to the training dataset. Different noise\_std were applied.

```
def split_data(df): 1 usage
    train_ratio = 0.6
    val_ratio = 0.2
    test_ratio = 0.2
    random_state = 42
    df = df.copy()
    train_val, test = train_test_split(*arrays: df, test_size=test_ratio, random_state=random_state)
    train_val, test = train_test_split(*arrays: train_val, test_size=val_ratio/(train_ratio + val_ratio),

    return train_val, test
```

```
#dataset splitting
train_val, test = split_data(D1)
```

```
#add noise to the training dataset
train_noise = noisify(train)

#applying scaler to sub-sets
train_std, val_std, test_std = standardize_data(train_noise, val, test)
```

```
def noisify(df): 1 usage
    noise_std_dict = {
        'bill_length_mm': 3.0,
        'bill_depth_mm': 1.5,
        'flipper_length_mm': 7.0,
        'body_mass_g': 150.0
    }

    df = df.copy()

    for col, std in noise_std_dict.items():
        if col in df.columns:
            noise = np.random.normal(loc=0, scale=std, size=df[col].shape)
            df[col] += noise
        else:
            print(f"Column '{col}' not found in the dataset.")
```



# Feature extraction

We extracted the sub-sets to excel files as requested.

[cmpsML\\_Tech19\CODE\OUTPUT\feature\\_extracted\\_data.xlsx](#)

[cmpsML\\_Tech19\CODE\INPUT\TEST\val\\_test\\_std.xlsx](#)

[cmpsML\\_Tech19\CODE\INPUT\TRAIN\train\\_std.xlsx](#)

Discussion:

```
def standardize_data(df1,df2,df3): 1 usage
    df1 = df1.copy()
    df2 = df2.copy()
    df3 = df3.copy()
    numeric_cols = ['bill_length_mm', 'bill_depth_mm', 'flipper_length_mm', 'body_mass_g']
    scaler = StandardScaler()
    cols_to_scale = [col for col in numeric_cols if col in df1.columns]
    df1[cols_to_scale] = scaler.fit_transform(df1[cols_to_scale])
    df2[cols_to_scale] = scaler.transform(df2[cols_to_scale])
    df3[cols_to_scale] = scaler.transform(df3[cols_to_scale])
    return df1,df2,df3
```

In PA2 we made a mistake that applying std-scaler **before** splitting the dataset.

This will lead to **data leakage**. When we scale based on the entire dataset, the scaling parameters (mean and standard deviation) are influenced by the validation and test data. This means that **information from the validation and test sets leaks into the training process**. Also we should only apply ‘**fit\_transform**’ to training dataset, and ‘**transform**’ to val and test datasets.

We fixed this issue after discovering it. Split the dataset into train, val and test first then scale the sub-sets separately.

# Description of the feature data

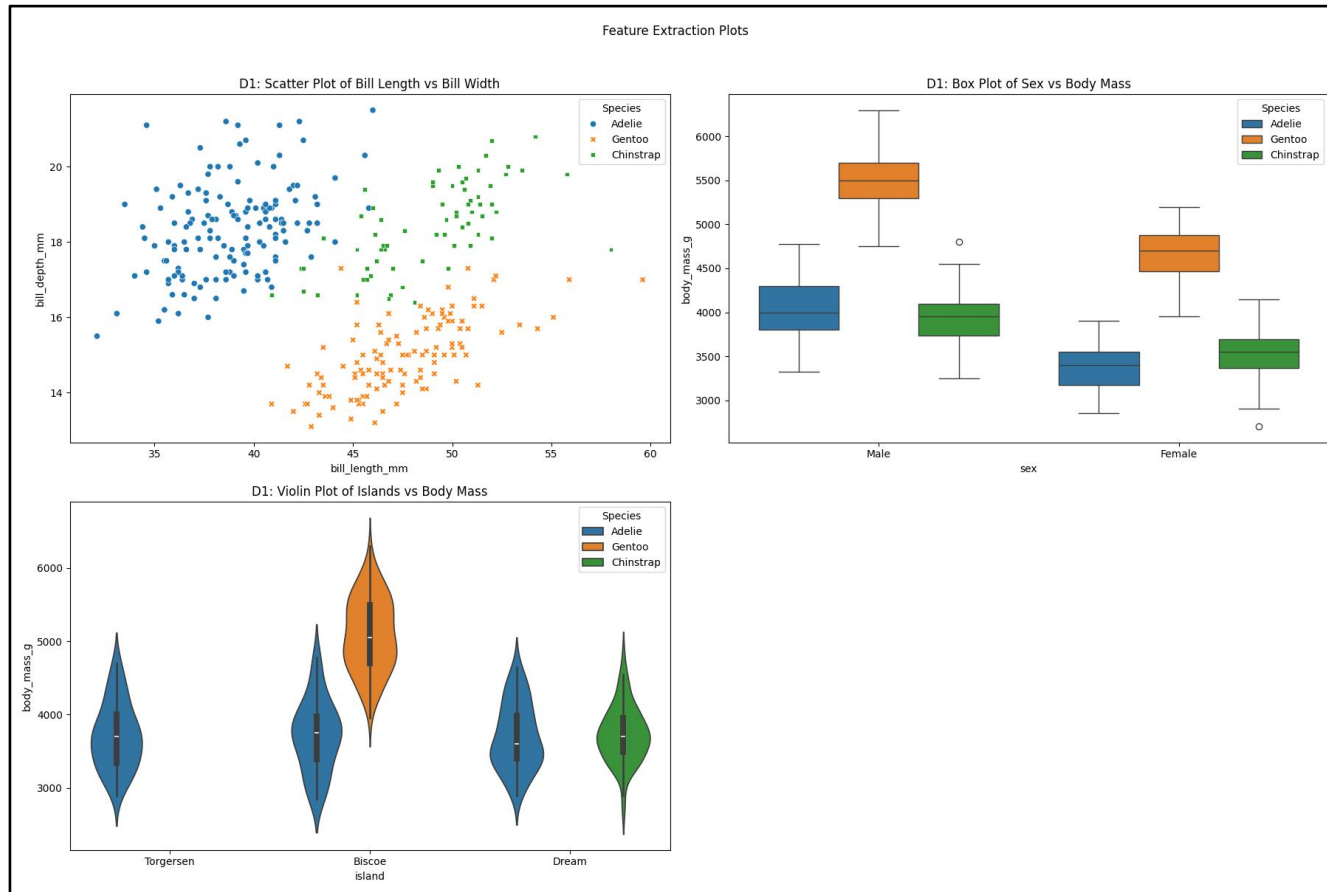
We have 9 numerical columns after feature extraction:

Features	Type	Notes
species	Numerical	Target, label encoding, 0 = 'Adelie', 1 = 'Chinstrap', 2 = 'Gentoo'
bill_length_mm	Numerical	Standardized
bill_depth_mm	Numerical	Standardized
flipper_length_mm	Numerical	Standardized
body_mass_g	Numerical	Standardized
sex_male	Numerical	one-hot encoding, 1 = 'male', 0 = 'female'
island_Biscoe	Numerical	one-hot encoding, 1 = True, 0 = False
island_Dream	Numerical	one-hot encoding, 1 = True, 0 = False
island_Torgersen	Numerical	one-hot encoding, 1 = True, 0 = False

[cmpsML\\_Tech19\OTHER\data\\_description.xlsx\feature\\_data\\_description](#)

# Description of the feature data

Here is a feature plot of original dataset:



Discussion:

From the first plot on top, left, we can see that these three penguin species **do not have much overlap** in bill length and width distribution, which means **they are separable in this feature**. (This is why we added some **noise** later.)

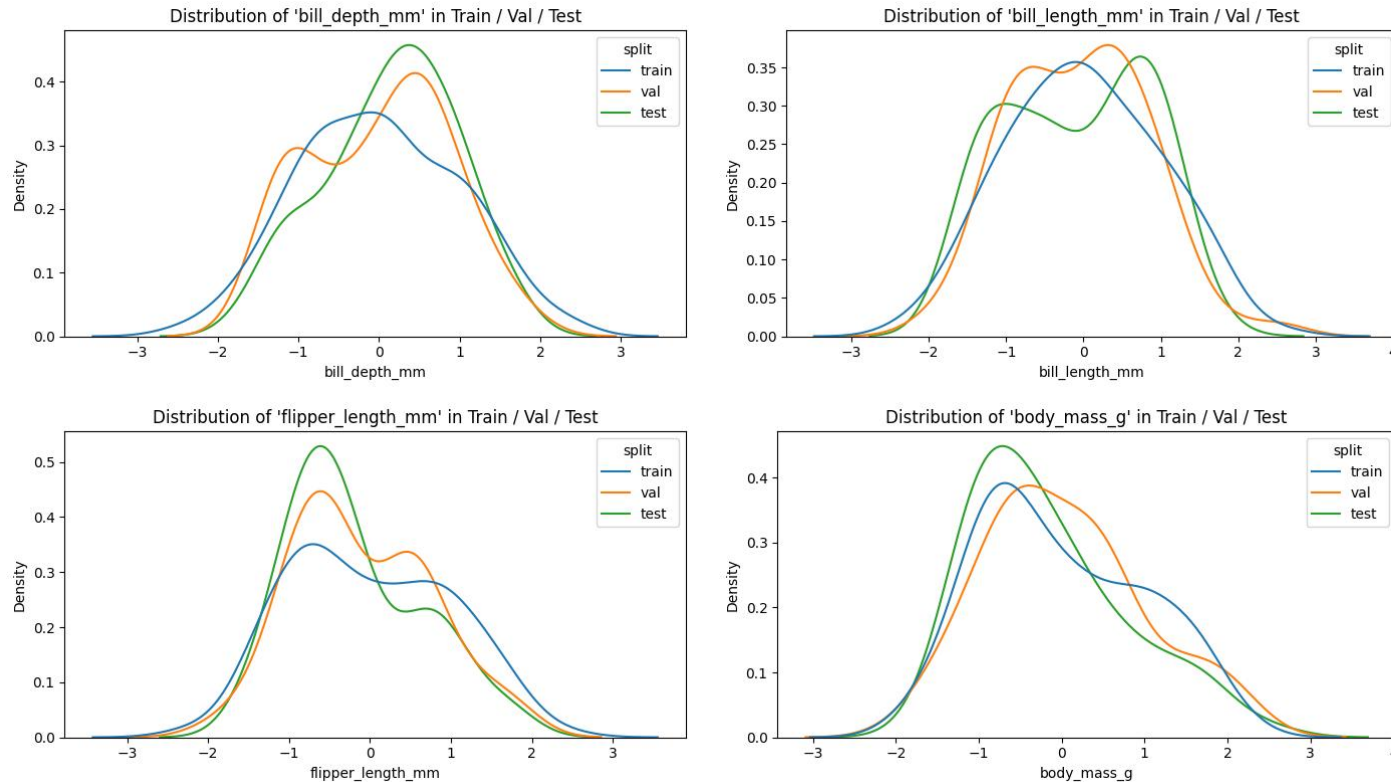
From the second plot on top, right, we can see that **Gentoo is much heavier** than the other two species.

From the third plot on bottom, left, we can see that Gentoo only lives on island Biscoe. Chinstrap only lives on island Dream. On island Torgersen, there is only one specie - Adelie. This makes **island and body mass good separation nodes in DT**.

The feature island highly correlates to species. **We decided to drop the 'island' column during the training.**

# Description of the feature data

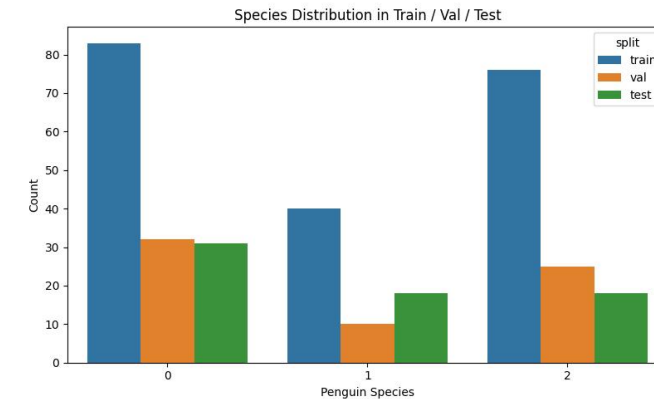
Here are our distribution of train/val/test datasets.



Discussion:

The distribution looks similar in val/test datasets, and a little bit different in training dataset due to the noise we added.

Basically our operation is reasonable.



# Description of the model

We built 5 models: KNN, DT, SVM, Grid search ANN and Random setting ANN.

```
'KNN': KNeighborsClassifier(  
    n_neighbors=5,  
    weights='distance',  
    p=2  
)  
  
'DT': DecisionTreeClassifier(  
    criterion='gini',  
    max_depth=4,  
    min_samples_split=5,  
    random_state=42  
)  
  
'SVM': SVC(  
    C=1.0,  
    kernel='rbf',  
    gamma='scale',  
    probability=True,  
    random_state=42  
)
```

Model parameters description

KNN:

n\_neighbours, k = 5

weights, weight function for neighbours (uniform/distance)

p, 1 = L1 distance, 2 = L2 distance

DT:

criterion, function to measure the quality of a split (gini/entropy)

min\_samples\_split, minimum number of samples required to split

SVM:

c, regularization parameter, default setting is 1.0.

# Description of the model

```
'RANDOM_ANN': MLPClassifier(  
    solver='adam',  
    batch_size=60,  
    max_iter=50,  
    early_stopping=True,  
    random_state=42,  
    hidden_layer_sizes=(6,3),  
    activation = 'relu',  
    alpha = 0.001,  
    learning_rate_init = 0.05,  
    validation_fraction=0.2  
)
```

Model parameters description

ANN:

solver, the solver for weight optimization

batch\_size = 60, size of minibatches for optimizers

max\_iter, maximum number of iterations

early\_stopping = True, stop when validation score is not improving

hidden\_layer\_size = (6,3) , 6 nodes at first layer and 3 nodes at second layer

activation, the activation function of the network

alpha, L2 regularization

learning\_rate\_init, the initial learning rate

validation\_fraction, the proportion of training data to set aside as validation set for early stopping

# Description of the model

```
def grid_search_ann(X, y): 1 usage
    print("\nRunning GridSearch for ANN...")
    # define pipeline and hyperparameter grid
    ann_pipeline = Pipeline([
        ("scaler", StandardScaler()),
        ("clf", MLPClassifier(
            solver='adam',
            batch_size='auto',
            max_iter=300,
            early_stopping=False,
            random_state=42
        ))
    ])

    param_grid = {
        "clf__hidden_layer_sizes": [(8, 4), (6, 3), (4, 2)],
        "clf__activation": ['relu', 'tanh'],
        "clf__alpha": [0.0001, 0.001],
        "clf__learning_rate_init": [0.001, 0.005, 0.01]
    }

    # Stratified K-Fold cross-validation
    skf = StratifiedKFold(n_splits=3, shuffle=True, random_state=42)
    grid_search = GridSearchCV(
        estimator=ann_pipeline,
        param_grid=param_grid,
        cv=skf,
        scoring='accuracy',
        n_jobs=-1,
        verbose=1
    )
    grid_search.fit(X, y)
```

Grid search was applied to the grid\_search\_ann.

3-fold cross validation was also applied.

The rest models used 3-tier testing scheme.

We intentionally set the ANN to be a small one.


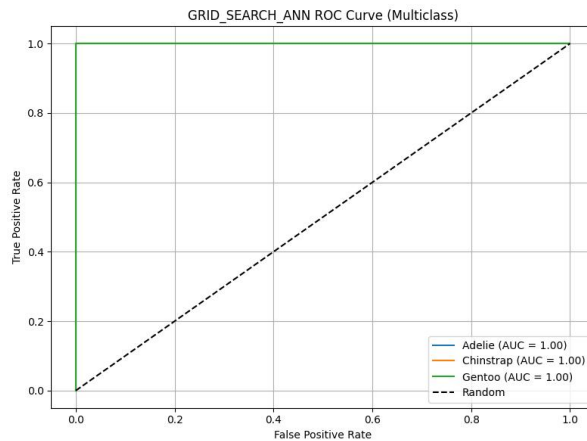
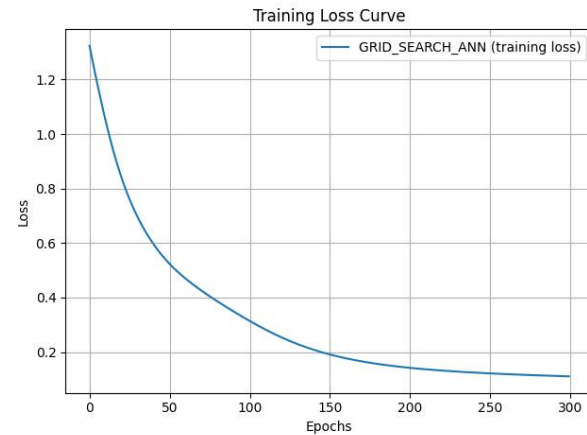
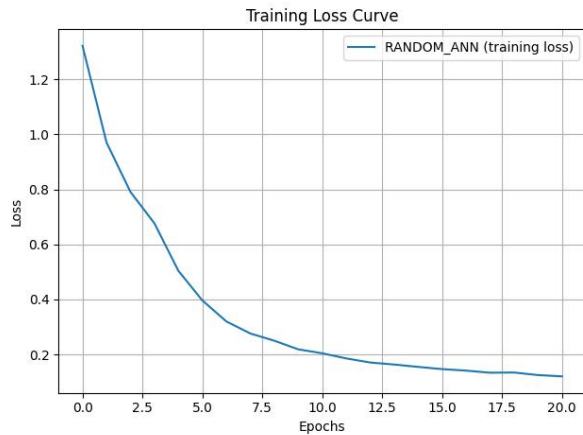


# Performance of the model

All the plots and excel files are under the folder:

`cmpsML_Tech19\CODE\OUTPUT\model performance`

Grid search K-fold cross validation results:



A	B
Fold	Accuracy
1	0.955224
2	0.939394
3	0.939394
Mean	0.944671

Training loss curve

Roc curve

for ANN

(MLPClassifier doesn't provide epoch-error curve for val)

# Performance of the model

## Performance plots for KNN

### KNN Confusion Matrix:

	Adelie	Chinstrap	Gentoo
Adelie	31	0	0
Chinstrap	1	17	0
Gentoo	0	0	18

KNN Average ROC AUC: 1.0000

### KNN Evaluation Metrics:

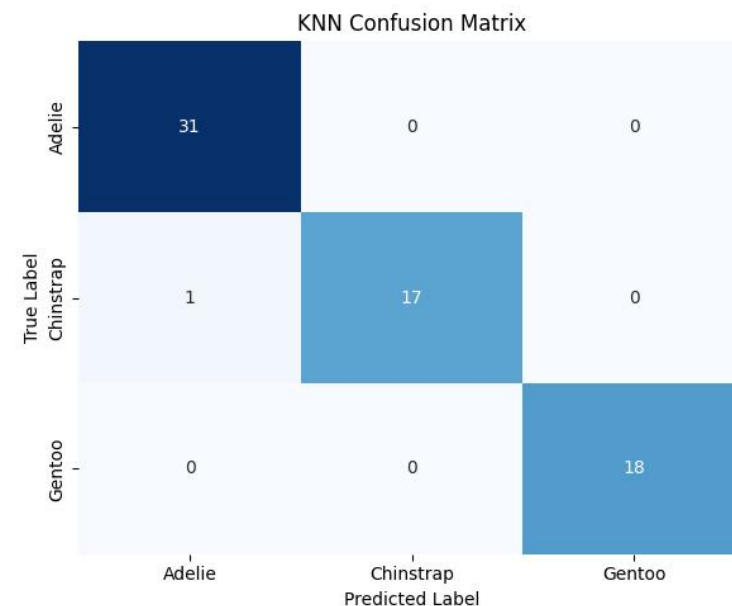
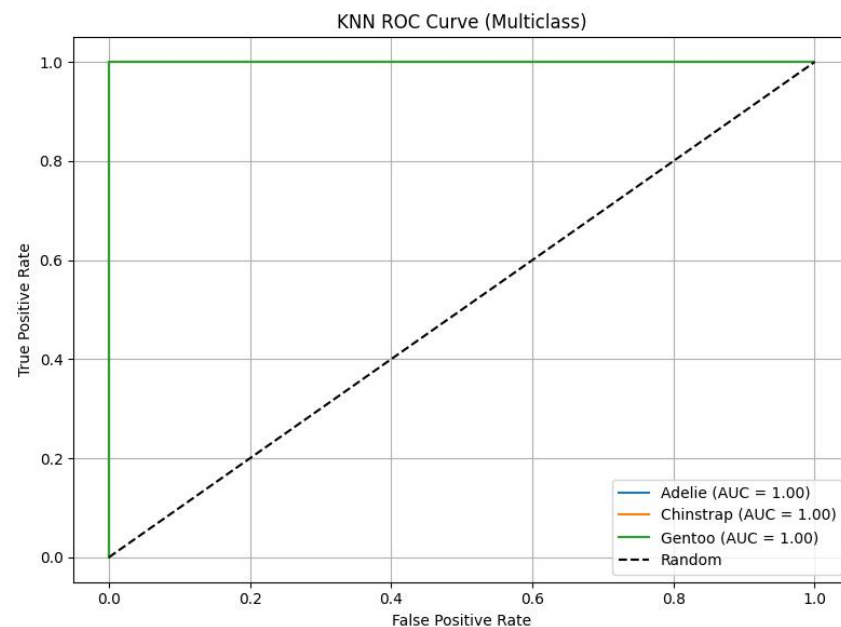
Accuracy: 0.9851

Precision: 0.9896

Recall: 0.9815

Specificity (avg): 0.9907

F1 Score: 0.9852



# Performance of the model

## Performance plots for SVM

### SVM Confusion Matrix:

	Adelie	Chinstrap	Gentoo
Adelie	31	0	0
Chinstrap	2	16	0
Gentoo	0	0	18

SVM Average ROC AUC: 1.0000

### SVM Evaluation Metrics:

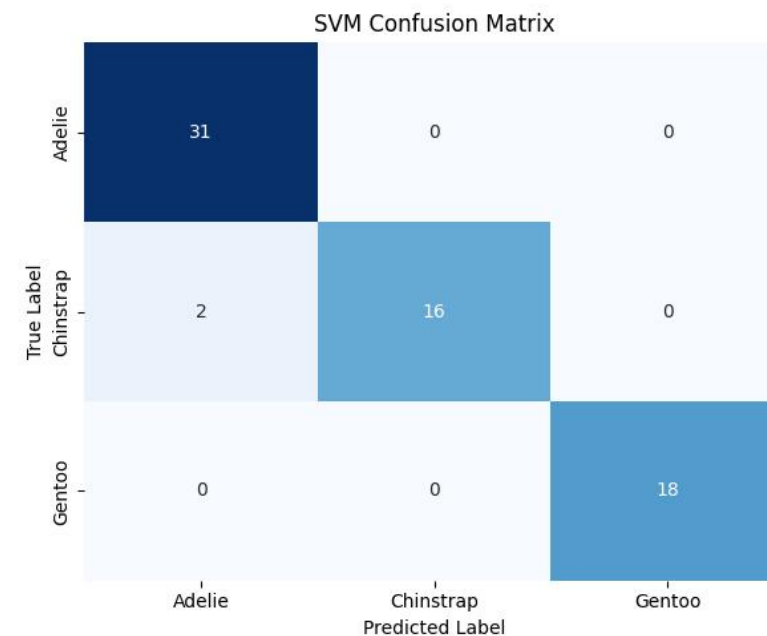
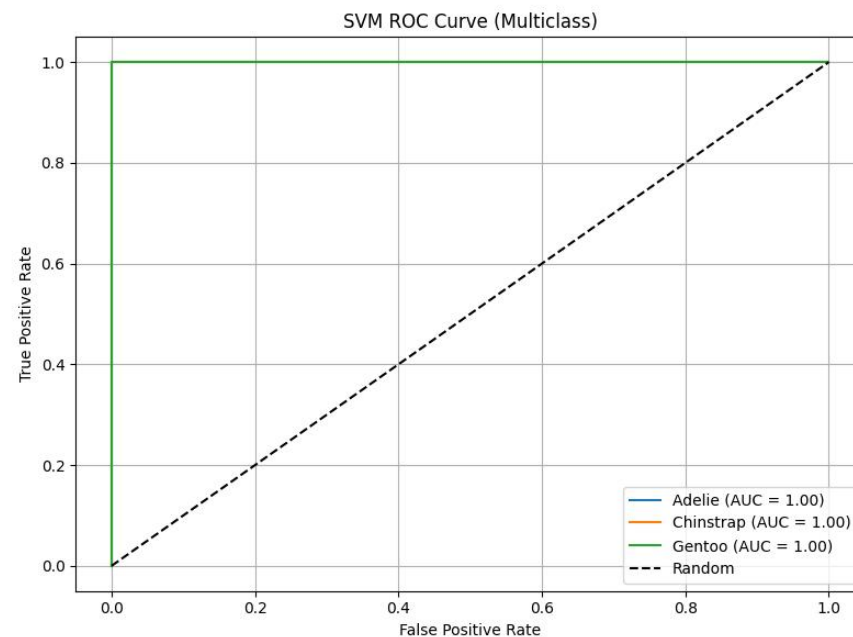
Accuracy: 0.9701

Precision: 0.9798

Recall: 0.9630

Specificity (avg): 0.9815

F1 Score: 0.9700



# Performance of the model

## Performance plots for DT

### DT Confusion Matrix:

	Adelie	Chinstrap	Gentoo
Adelie	31	0	0
Chinstrap	2	16	0
Gentoo	0	0	18

DT Average ROC AUC: 0.9946

### DT Evaluation Metrics:

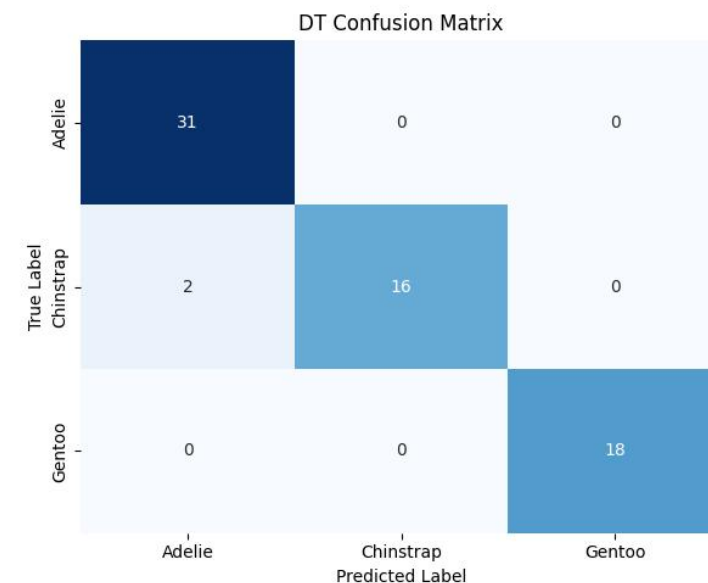
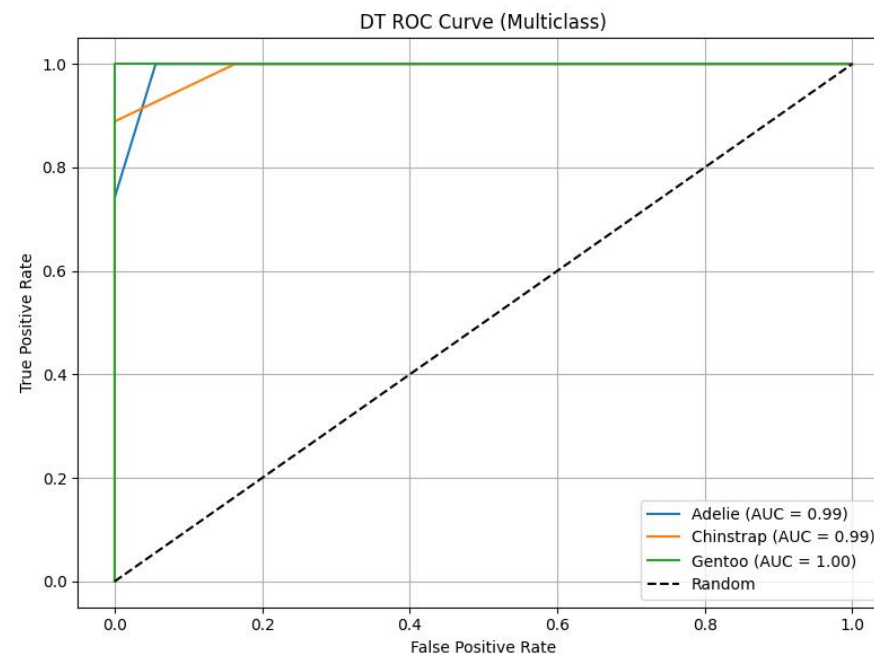
Accuracy: 0.9701

Precision: 0.9798

Recall: 0.9630

Specificity (avg): 0.9815

F1 Score: 0.9700



# Performance of the model

## Performance plots for RANDOM\_ANN

RANDOM\_ANN Confusion Matrix:

	Adelie	Chinstrap	Gentoo
Adelie	31	0	0
Chinstrap	2	16	0
Gentoo	0	0	18

RANDOM\_ANN Average ROC AUC: 1.0000

RANDOM\_ANN Evaluation Metrics:

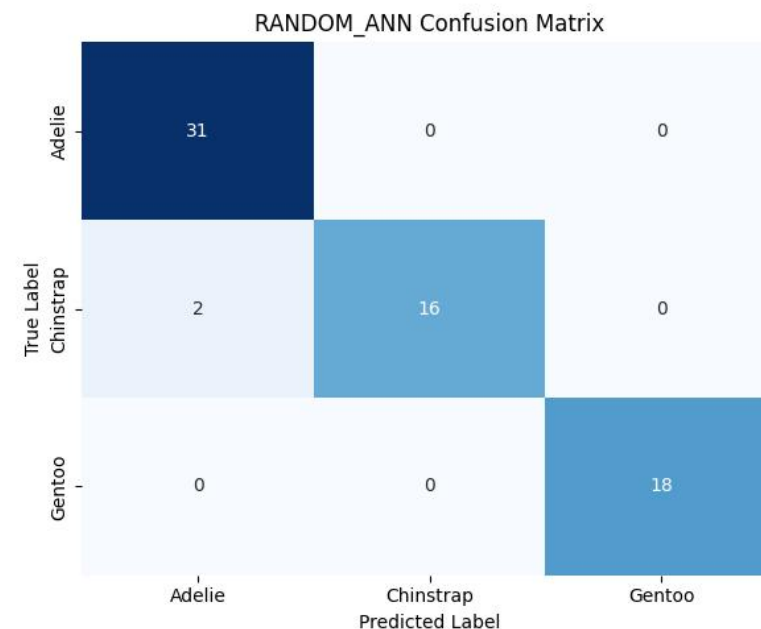
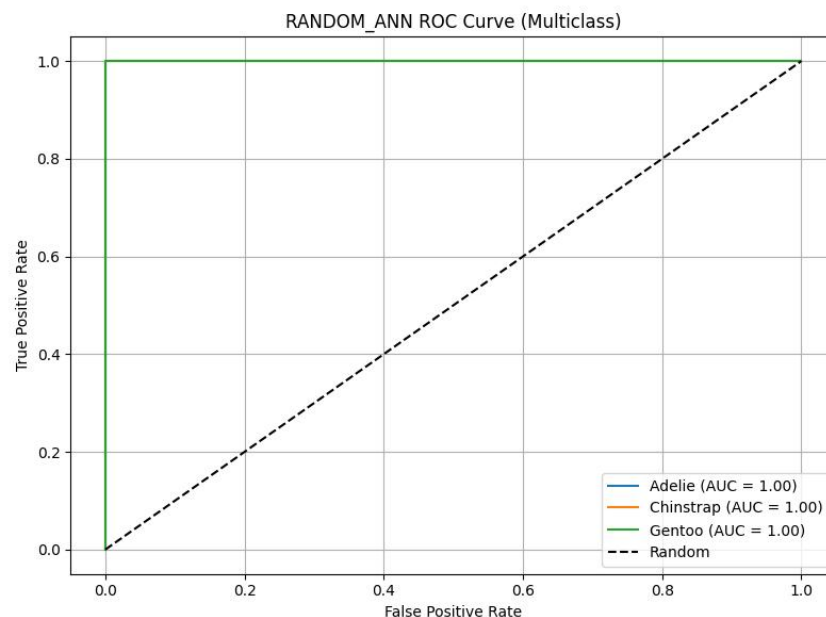
Accuracy: 0.9701

Precision: 0.9798

Recall: 0.9630

Specificity (avg): 0.9815

F1 Score: 0.9700



# Performance of the model

## Performance plots for GRID\_SEARCH\_ANN



GRID\_SEARCH\_ANN Confusion Matrix:

	Adelie	Chinstrap	Gentoo
Adelie	31	0	0
Chinstrap	2	16	0
Gentoo	0	0	18

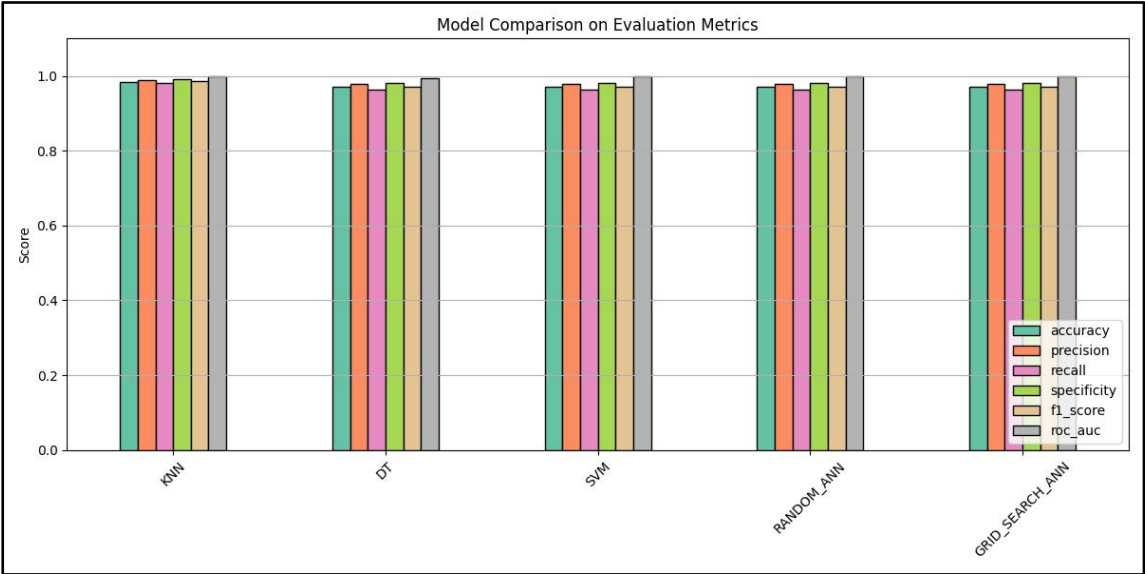
GRID\_SEARCH\_ANN Average ROC AUC: 1.0000

GRID\_SEARCH\_ANN Evaluation Metrics:

Accuracy: 0.9701  
Precision: 0.9798  
Recall: 0.9630  
Specificity (avg): 0.9815  
F1 Score: 0.9700



# Performance of the model



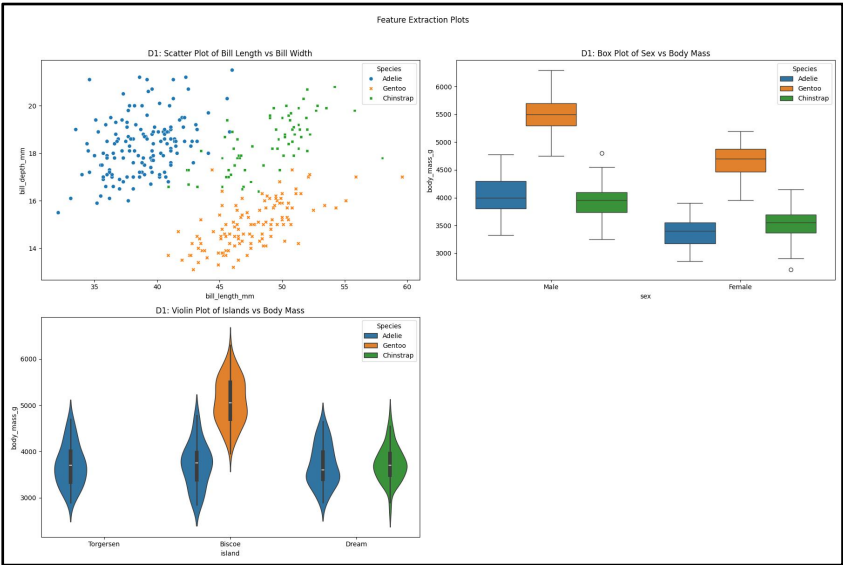
Discussion:

We found that all models are working very well with this penguins classification task.

We have been very careful to avoid data leakage, and tried several ways to reduce the accuracy:

Add noise to the dataset,  
limit the nodes and layers of ANN,  
drop the island column of the dataset.

Still, the penguins are separable in some feature, making the task easy to complete.

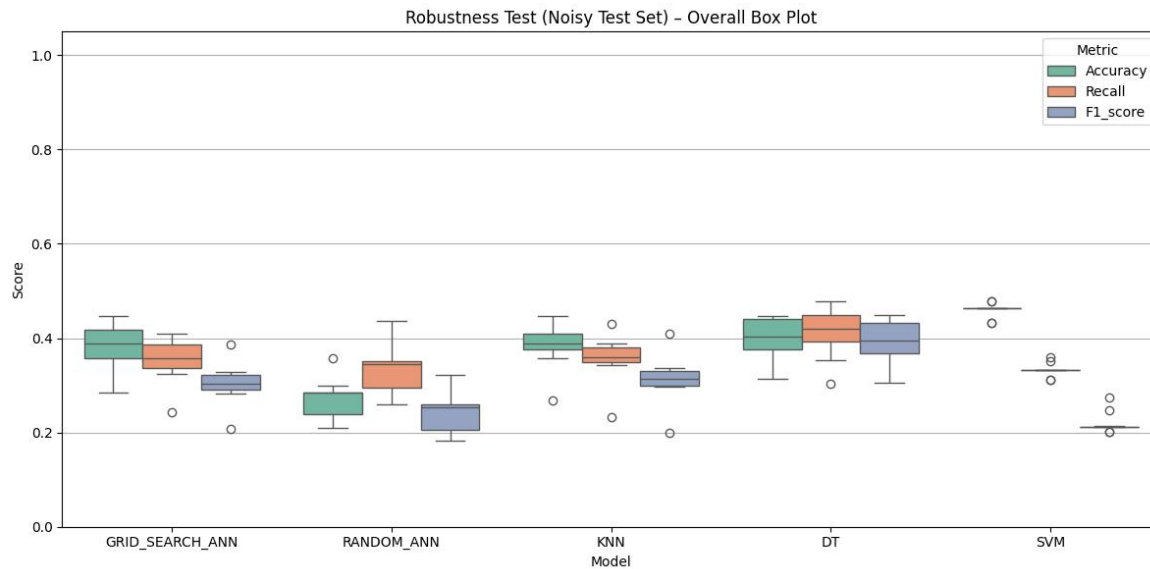


	A	B	C	D	E	F	G
		accuracy	precision	recall	specificity	f1_score	roc_auc
	KNN	0.985074627	0.989583333	0.981481481	0.990740741	0.985185185	1
	DT	0.970149254	0.97979798	0.962962963	0.981481481	0.96997549	0.994587
	SVM	0.970149254	0.97979798	0.962962963	0.981481481	0.96997549	1
	RANDOM_ANN	0.970149254	0.97979798	0.962962963	0.981481481	0.96997549	1
	GRID_SEARCH_ANN	0.970149254	0.97979798	0.962962963	0.981481481	0.96997549	1



# Performance of the model

## Robustness test



```
def noisify(df):  
    1 usage  
    noise_std_dict = {  
        'bill_length_mm': 3.0,  
        'bill_depth_mm': 1.5,  
        'flipper_length_mm': 7.0,  
        'body_mass_g': 150.0  
    }
```

## Discussion:

We conducted 10 trials to test the robustness of the models.

The result decreased significantly because we added strong noise to the test dataset.

# Team members & roles / Task completion report

Team members: Drew Hutchinson, Zichuo Wang

MS only tasks were done by Zichuo Wang.

Besides, all the work was accomplished through our discussion, from proposal to the final report writting.

A	B	C	D
Date	Task Name	Status	Person
3/14/2025	Proposal	Done	Drew Hutchinson, Zichuo Wang
3/31/2025	PA1	Done	Drew Hutchinson, Zichuo Wang
4/14/2025	PA2	Done	Drew Hutchinson, Zichuo Wang
5/8/2025	MS only part	Done	Zichuo Wang
5/8/2025	Coding and fixing	Done	Drew Hutchinson, Zichuo Wang
5/8/2025	Final report writting	Done	Zichuo Wang
5/8/2025	Review and check	Done	Drew Hutchinson